

GIT BUILD SERIES

ANSWER SHEET

Workbook 2: Skills, Workflows & GitHub Repos

Complete answer key for all 8 module quizzes, exercise suggested responses, and the 20-question final assessment. For instructor use and self-study verification.

HOW TO USE

Answer Sheet Guide

This answer sheet provides the correct responses for all quizzes and suggested approaches for all exercises in GIT BUILD Workbook 2. Quiz answers include the correct option letter plus a brief explanation of why it is correct. Exercise responses represent model answers that demonstrate the expected depth and quality of student work. Instructors should use these as rubric guidelines rather than requiring exact verbatim matches, as students may approach exercises from different but equally valid angles. The final assessment scoring guide at the end will help you evaluate overall comprehension across all eight modules.

NOTE: Exercise answers are suggested responses, not the only correct answers. Students who demonstrate critical thinking, real-world application, and a clear understanding of the underlying concepts should receive full credit even if their specific answers differ from these models.

Quiz Answers: Coding Agents & IDEs

1. Which coding agent is Anthropic's official terminal-based tool?

B) Claude Code

Claude Code is Anthropic's official CLI-based coding agent that operates directly within your repository, providing terminal-native AI-assisted development.

2. What is the primary advantage of Cursor's "Composer" feature?

C) It creates multi-file implementations from plain English descriptions

Composer goes beyond single-line completions by understanding cross-file dependencies and generating coordinated changes across multiple files simultaneously from natural language prompts.

3. Which tool is best described as an open-source AI pair programmer with deep Git integration?

C) Aider

Aider is an open-source terminal-based AI pair programmer that integrates deeply with Git, automatically committing changes and managing branches as it works alongside you.

4. What development philosophy does Spec Kit enforce?

B) Spec-first workflow

Spec Kit enforces a spec-first approach where you define what you want to build before writing code, significantly reducing iteration cycles by establishing clear requirements upfront.

5. If a freelancer delivers a 20-hour project in 4 hours at \$100/hour, what is their effective hourly rate?

D) \$500

$100/\text{hour} \times 20\text{hours} = 2,000$ total project value. $2,000/4\text{actualhours} = 500$ effective hourly rate. This demonstrates how AI coding agents multiply earning potential.

EXERCISE: SET UP YOUR AGENT STACK — SUGGESTED RESPONSE

1. **Primary agent choice:** Claude Code is recommended for terminal-native workflows, while Cursor is better for developers who prefer a GUI. Claude Code integrates directly with your repo and supports complex multi-step tasks. Document the installation process and any initial setup steps required.
2. **Test project structure:** A typical test project should include a README.md with project description, a main.py or index.js entry point, a requirements.txt or package.json for dependencies, and a .gitignore. Running git init and making an initial commit establishes version control.
3. **Agent-built application:** When asking the agent to build a to-do API, a strong prompt specifies the framework (e.g., Python/Flask), endpoints (GET/POST/PUT/DELETE for tasks), data model fields (id, title, completed, created_at), and any middleware (CORS, error handling). Most agents will generate a working API in under two minutes.
4. **Review findings:** Common observations include: agents excel at boilerplate and standard patterns but may struggle with business logic nuances. Iteration count typically ranges from 2-5 for a complete, polished result. Testing generation quality varies by agent.
5. **Time comparison:** A typical to-do API that takes 2-4 hours manually can often be completed in 15-30 minutes with AI assistance, representing a 4-10x productivity improvement for standard tasks.

Quiz Answers: Agent Frameworks

1. Which framework uses directed graphs for agent workflows?

B) LangGraph

LangGraph uses directed graphs to define agent workflows, providing fine-grained control over state transitions, decision points, and complex multi-step reasoning chains.

2. What is CrewAI's primary value proposition?

B) It enables specialized agents with defined roles to collaborate on complex tasks

CrewAI's multi-agent orchestration lets you create teams of specialized agents (researcher, writer, reviewer) that collaborate like a real team, each contributing their expertise.

3. Which framework is best for non-technical users to design AI workflows visually?

C) Dify

Dify provides a visual drag-and-drop builder for designing AI workflows, making it accessible to non-technical team members who can prototype and deploy agents without writing code.

4. What makes pydantic-ai unique?

A) It uses Pydantic models for type-safe agent interactions

pydantic-ai leverages Pydantic's type system to enforce structured, type-safe interactions between agents and their tools, catching errors at development time rather than runtime.

5. An agent saving \$50,000/year could be priced at what monthly rate?

C) \$500-1,000/month

Standard SaaS pricing targets 10-20x the value delivered monthly. $50,000/year = 4,167/month$.

Charging \$500-1,000/month represents 12-24% of monthly value, a competitive and sustainable price point.

EXERCISE: DESIGN YOUR FIRST AI AGENT — SUGGESTED RESPONSE

1. **Problem identification:** A strong candidate problem is one that is repetitive, rule-based but with judgment calls, and currently done manually. Examples: processing customer support emails, generating weekly status reports from Jira tickets, or qualifying sales leads based on criteria. The problem statement should include who does it, frequency, time per occurrence, and current pain points.
2. **Framework choice with justification:** For a single-task automation, pydantic-ai provides clean type safety and minimal overhead. For multi-step research workflows, LangGraph's graph structure handles branching logic well. For a team simulation (e.g., content creation pipeline), CrewAI's role-based agents shine. Dify is ideal if non-technical stakeholders need to modify the workflow later.
3. **Architecture outline:** A typical agent architecture includes: Input node (receives raw data), Processing node (applies AI analysis), Decision node (routes based on classification), Action node (executes the appropriate response), and Output node (formats and delivers results). Each node should have defined inputs and outputs.
4. **Tool specifications:** Each tool needs: a descriptive name, a clear purpose statement, typed input parameters (use Pydantic BaseModel or JSON Schema), and a structured output format. Example: WebSearchTool(query: str, max_results: int) returns list of SearchResult objects with url, title, snippet fields.
5. **MVP plan:** Start with the core loop: receive input, process with AI, return output. Estimate 2-4 days for a working prototype with AI coding tools. Week 1: core functionality. Week 2: error handling and edge cases. Week 3: testing and user feedback. Week 4: polish and deployment.

Quiz Answers: MCP Servers & Tooling

1. What analogy best describes MCP's role?

B) A universal connector (like USB-C) for AI models and tools

MCP acts as a standardized protocol that lets any AI model connect to any tool or data source, similar to how USB-C provides a universal physical connector standard across devices.

2. Why is Tavily's API suited for AI agents?

B) It returns results optimized for LLM consumption rather than human browsing

Tavily structures search results specifically for AI agents: clean content, relevant snippets, and structured data formats that LLMs can process directly without additional parsing or extraction.

3. What problem does Context7 solve?

B) It provides up-to-date library documentation on demand instead of relying on outdated training data

Context7 fetches current documentation from libraries and frameworks in real-time, solving the common problem of AI agents relying on outdated or hallucinated API information from their training data.

4. Which library is for rapidly building custom MCP servers?

C) fastmcp

fastmcp provides a minimal, Pythonic API for defining MCP servers with decorators and type hints, allowing developers to create functional MCP servers in just a few lines of code.

5. At 49/month, how many customers store each 10K MRR?

C) 204

$10,000 / 49 = 204.08$, so approximately 204 paying customers. This is a realistic target for a well-built MCP server serving a professional developer audience.

EXERCISE: BUILD A CUSTOM MCP SERVER — SUGGESTED RESPONSE

1. **Tool gap identification:** Good candidates include: a server that connects to Notion databases, one that queries Shopify store analytics, one that interfaces with local government data APIs, or one that parses and extracts data from PDF invoices. The gap should be specific enough that a focused MCP server solves it better than a general-purpose tool.
2. **Tool definitions:** Example for a "GitHub Analytics MCP Server": (1) `get_repo_stats(owner, repo)` returns stars, forks, issues, PR stats; (2) `get_contributor_activity(owner, repo, days)` returns commit frequency per contributor; (3) `get_issue_timeline(owner, repo)` returns issue creation and resolution trends. Each tool should have typed parameters and structured JSON output.
3. **Build approach:** `fastmcp` is recommended for most use cases due to its minimal boilerplate and Python ecosystem. Use the raw MCP SDK only if you need fine-grained control over transport layers or are building in TypeScript. Justification should reference development speed, maintainability, and community support.
4. **Specification document:** Should include: server name (e.g., "github-analytics-mcp"), description, version, transport method (stdio for Claude Code, SSE for web applications), list of tools with full JSON schemas for inputs/outputs, any resources (e.g., cached data), and prompts (e.g., "Analyze this repository's health").
5. **Go-to-market plan:** Recommended approach: open-source core with premium features (advanced analytics, team dashboards) at \$29-49/month. Launch on GitHub, list on MCPHub, create a demo video, post on Hacker News and Reddit r/ClaudeAI. Target first 50 users through community engagement within 30 days.

Quiz Answers: Local AI & Infrastructure

1. What command downloads and runs Llama 3.1 with Ollama?

B) ollama pull llama3.1 then ollama run llama3.1

Ollama uses a two-step process: "pull" downloads the model weights, and "run" starts an interactive chat session. This separation lets you download models ahead of time and run them offline.

2. What is LlamaFile's key innovation?

B) It packages a model into a single executable file that runs without installation

LlamaFile (by Mozilla) combines the model weights and inference engine into a single executable that runs on any operating system without separate installation steps, dramatically simplifying AI model distribution.

3. How does Unsloth achieve faster fine-tuning?

B) Through custom Triton kernels, optimized attention, and smart memory management

Unsloth uses hand-optimized Triton kernels for GPU computation, implements efficient attention mechanisms, and manages GPU memory strategically, achieving 2-5x speedups with 80% less VRAM usage.

4. What technique does vLLM use for faster inference?

B) PagedAttention and continuous batching

vLLM's PagedAttention manages KV cache memory efficiently (similar to virtual memory paging), while continuous batching eliminates idle GPU time between requests, together achieving 5-10x throughput improvement.

5. A SaaS processing 1M tokens/day at \$15/day API cost would save how much per year switching to local?

C) \$5,400

$15/\text{day} \times 365 \text{ days} = 5,475/\text{year}$ in API costs. With local inference on owned hardware (electricity only), this becomes near-zero marginal cost, saving approximately \$5,400/year. The hardware investment typically pays for itself within 3-6 months at this scale.

EXERCISE: DEPLOY A LOCAL AI STACK — SUGGESTED RESPONSE

1. **Ollama installation:** After downloading from ollama.ai and installing, verify with "ollama list" which should return an empty list initially. Common issues include PATH configuration on Linux (add to .bashrc) and firewall permissions on corporate networks.
2. **Model testing:** After "ollama pull llama3.1" (approximately 4.7GB download), run "ollama run llama3.1" for an interactive session. Llama 3.1 8B performs well on general knowledge and coding tasks but may lag behind Claude/GPT-4 on complex reasoning. Test with both factual questions and creative prompts to gauge range.
3. **Open WebUI setup:** Open WebUI provides a ChatGPT-like interface for local models. After installation via Docker (recommended), connect it to Ollama's API at localhost:11434. Test chat functionality, PDF/file upload for RAG-style queries, and model switching between multiple downloaded models.
4. **Application connection:** Ollama exposes an OpenAI-compatible API at localhost:11434/v1. A simple Python script using the openai library with base_url="http://localhost:11434/v1" can connect directly. Alternatively, configure Claude Code or Aider to use Ollama as a backend provider.
5. **Savings calculation example:** If using 10M tokens/month at Claude's $3 / M_{input} + 15 / M_{output}$ rate (blended $\sim 8 / M_{average}$), $monthlycost is 80$, or $960 / year$. $Local inference electricity cost for a consumer GPU running 8 hours / day is approximately 5-15 / month (60 - 180 / year)$, representing 780-900 in annual savings.

Quiz Answers: Design & Creative Skills

1. Why is design the highest-demand AI skill category in 2025?

B) Every business needs design, most can't afford designers, and AI can produce 80-90% quality output

The combination of universal demand (every business needs visual assets), high cost of human designers (\$75-200/hour), and AI's ability to produce near-professional quality makes design skills the most monetizable AI skill category.

2. What makes Hand-Drawn Diagrams uniquely valuable?

B) Its sketch-like, approachable aesthetic feels human and is popular in startups and education

Hand-drawn diagrams bypass the "too polished" feeling of standard AI output. The Excalidraw-style aesthetic feels approachable and collaborative, making it particularly effective for presentations, educational content, and startup pitches.

3. Which design skill category commands the highest per-project rates?

C) 3D/Motion — requires most creative direction and produces most impressive outputs

3D and motion design using Three.js, GSAP, and React Three Fiber command \$1,000-10,000 per project because the technical barrier is higher, the output is more visually impressive, and there are fewer practitioners who can combine AI tools with these technologies.

4. What is the strategic insight about AI and the design agency model?

B) A single skilled operator with AI tools can produce agency-quality output, replacing large teams

The key insight is that AI tools compress the design production pipeline. Where an agency needed 5-10 specialists (strategist, designer, developer, project manager), one skilled person with AI tools can produce comparable output, fundamentally disrupting the agency economics.

5. What accessibility standard does Color Expert ensure?

C) WCAG contrast requirements

The Color Expert skill ensures all color combinations meet WCAG (Web Content Accessibility Guidelines) contrast ratios, making designs accessible to users with visual impairments and complying with legal requirements in many jurisdictions.

EXERCISE: BUILD A DESIGN PORTFOLIO IN 30 MINUTES — SUGGESTED RESPONSE

1. **Workspace setup:** Create the project folder and ensure Claude Code or your AI tool is running with the frontend-design skill loaded. Set up a simple structure: /designs/landing.html, /designs/portfolio.html, /designs/product.html for the three outputs.
2. **Design #1 (SaaS Landing):** A good prompt specifies: "Build a SaaS landing page for 'FlowMetrics' — a project management tool for remote teams. Minimal style with navy and white. Include hero, 3 feature cards, pricing table, and CTA footer." Expected time: 5-8 minutes including iteration.
3. **Design #2 (Portfolio):** Try a contrasting style: "Build a portfolio for 'Alex Chen', a street photographer. Bold, editorial style with black/white/one accent color. Grid gallery with hover effects, about section, and contact form." Expected time: 5-8 minutes.
4. **Design #3 (E-commerce):** "Design a product page for a premium mechanical keyboard. Warm, premium aesthetic with wood tones. Include image gallery, specs table, reviews, and related products carousel." Expected time: 5-8 minutes.
5. **Critique findings:** AI-generated designs typically score 7-9/10 on visual appeal, 6-8/10 on usability, and 7-8/10 on commercial viability. Common strengths: consistent spacing, modern aesthetics, responsive layouts. Common weaknesses: generic imagery placeholders, cookie-cutter layouts, lack of brand personality. Total time: 20-30 minutes vs. 20-40 hours manually.

Quiz Answers: Marketing & Content Skills

1. What is the estimated engagement increase from an optimized hook?

C) 300-500%

Research shows that optimized hooks using frameworks like PAS, AIDA, and curiosity gaps can increase initial engagement by 300-500% compared to generic openings, because the first 3 seconds determine whether someone stops scrolling.

2. What does GEO stand for and why is it important?

B) Generative Engine Optimization — it optimizes for AI-powered answer engines like ChatGPT and Perplexity

GEO is the evolution of SEO for the AI search era. As more people get answers from ChatGPT, Perplexity, and Google AI Overviews, content must be optimized for citation and inclusion in AI-generated responses, not just traditional link clicks.

3. What is the average ROI for email marketing?

C) 36forevery1 spent

Industry data consistently shows email marketing delivers approximately 36inrevenueforevery1 spent, making it one of the highest-ROI marketing channels available. The Email Marketing Bible skill teaches this entire system.

4. Which two factors are the top determinants of YouTube performance?

C) Thumbnail and title

YouTube's own data confirms that thumbnail and title are the two most impactful factors for video performance. They determine CTR (click-through rate), which is YouTube's primary ranking signal for recommending content to new viewers.

5. How much do agencies charge for competitive ad analysis?

C) \$2,000-10,000

Professional competitive ad analysis reports from agencies typically cost \$2,000-10,000 depending on depth, number of competitors analyzed, and the format of deliverables. The Competitive Ads Extractor skill automates this at near-zero marginal cost.

EXERCISE: CREATE A CONTENT MACHINE — SUGGESTED RESPONSE

1. **Content persona example:** "Startup Founder Sarah" — 32-year-old SaaS founder, technical background, struggles with marketing, spends time on Twitter/X and LinkedIn, wants actionable growth advice, reads long-form content but engages with short-form video. Pain points: limited budget, wearing too many hats, needs systems not one-off tactics.
2. **Three hooks for the same idea (e.g., "AI tools save time"):** Curiosity: "I replaced 3 tools with one AI setup. Here's what broke." Contrarian: "Stop buying AI tools. You only need these 4." Data: "73% of AI tool users are using them wrong. Are you one of them?" Each uses a different psychological trigger to capture attention.
3. **Post scoring and iteration:** A well-scored post rates 8+/10 on attention capture (strong hook), value delivery (specific, actionable tips), readability (short paragraphs, one idea per line), CTA (clear next step), and shareability (includes a reframable insight). First drafts typically score 5-6; iteration brings them to 7-9.
4. **Week content plan:** Monday: Twitter thread (educational), Tuesday: LinkedIn post (thought leadership), Wednesday: Reel/TikTok (visual demo), Thursday: Newsletter deep-dive (long-form), Friday: Twitter poll + engagement post. Each post should have a distinct format and complementary topic.
5. **Pricing reflection:** A content creation service using this system could reasonably charge \$2,000-5,000/month for a full content calendar with posts, hooks, and analytics. Key insight: the value isn't in the individual posts but in the systematic, consistent output that compounds over time.

Quiz Answers: Research & Intelligence

1. What makes the Deep Research Engine different from a single AI search?

B) It decomposes questions, searches multiple sources, evaluates credibility, and iteratively refines

The Deep Research Engine uses an 8-phase pipeline: question decomposition, multi-source search (Brave, Serper, Exa, Firecrawl), source credibility scoring, evidence extraction, contradiction detection, synthesis, and iterative refinement for depth.

2. Which skill acts as a quality control layer verifying claims?

C) Evidence-Based Dialogue

Evidence-Based Dialogue replaces sycophantic AI responses with structured critical analysis, requiring claims to be backed by sources with confidence levels. It operates in 5 modes including STEELMAN for debate prep and DECISION for tradeoff analysis.

3. What is a "knowledge moat"?

B) A depth of domain understanding built over time that competitors cannot easily replicate

In AI research, a knowledge moat refers to the compounding advantage of deep domain expertise that you build through repeated research cycles. While AI tools are accessible to everyone, the ability to ask better questions and evaluate answers critically creates an unreplicable competitive edge.

4. What time reduction can AI research tools achieve for market analysis?

B) 2 weeks reduced to 2 hours

A comprehensive market analysis that traditionally takes a junior analyst 2 weeks (80+ hours) of manual research across reports, filings, and news sources can be completed in approximately 2 hours using AI research tools, representing a 40x time reduction.

5. The Autoresearch Skill primarily differs because it:

B) Works autonomously over extended periods, iteratively deepening its understanding

Inspired by Karpathy's autoresearch concept, this skill operates autonomously over hours or days, continuously refining and expanding its research. It follows a "modify, verify, keep, repeat" loop that progressively deepens understanding without constant human guidance.

EXERCISE: RUN A DEEP RESEARCH SPRINT — SUGGESTED RESPONSE

1. **Research topic example:** "What are the most effective pricing strategies for AI-powered SaaS products targeting small businesses in 2025?" This topic is specific, commercially relevant, requires multiple data sources, and has real-world application value.
2. **Pre-research knowledge:** Document what you already know (e.g., freemium is common, value-based pricing is ideal). Questions to answer: What pricing models work best? What do customers actually pay? What are the churn implications? Credible sources: industry reports (OpenView, ProfitWell), competitor pricing pages, customer interviews, academic papers on SaaS economics.
3. **Research execution:** Start with broad queries across 3-4 search engines. Identify 8-12 high-quality sources including at least 2 data-driven reports, 3 competitor analyses, and 2-3 expert opinions. Extract specific data points (conversion rates, ARPU benchmarks, churn by pricing tier). Note contradictions (e.g., one source says freemium works, another says it doesn't).
4. **Synthesis report structure:** Executive Summary: "AI SaaS products targeting SMBs see highest LTV with usage-based pricing at \$49-199/month entry points, with freemium working only when the free tier has clear limitations." Key Findings: 3-5 specific, data-backed insights. Detailed Analysis: comparison of models, case studies, risk factors. Recommendations: 3-5 actionable next steps with expected outcomes.
5. **Evidence check:** Rate each finding 1-5. High-confidence findings (score 4-5) are backed by multiple sources or hard data. Low-confidence findings (score 1-2) need additional primary research (customer interviews, A/B tests) to validate. This honest assessment of certainty is what separates professional research from superficial summaries.

Quiz Answers: Product & Building

1. What is the core principle behind JTBD?

B) Discover what progress customers are trying to make in their lives, not what features they request

JTBD (Jobs To Be Done) focuses on understanding the underlying progress a customer is trying to make, rather than collecting feature requests. Customers "hire" products to get a job done, and understanding the job reveals the true product requirements.

2. What is the typical price range for an AI Transformation Discovery assessment?

C) \$5,000-50,000

AI Transformation Discovery assessments, which use frameworks like BCG's 10/20/70 and Deloitte's AI Maturity Model, typically command \$5,000-50,000 depending on organization size, depth of analysis, and the inclusion of implementation roadmaps.

3. What does Repomix do and why is it useful?

B) It packs entire codebases into a single file optimized for LLM consumption

Repomix consolidates an entire codebase into one AI-friendly file, preserving structure and context. This is invaluable when you need Claude or any LLM to understand, audit, or modify a complete project rather than individual files.

4. What pattern do successful AI-native entrepreneurs follow?

C) Identify pain points, build MVPs in 1-4 weeks with AI, launch fast, iterate based on feedback

The most successful AI-native entrepreneurs follow a rapid iteration cycle: find a real pain point, use AI tools to build an MVP in days rather than months, launch immediately, gather user feedback, and iterate. Speed of learning beats perfection of planning.

5. Which skill is "the single most valuable customer research skill for product builders"?

C) JTBD Interview Tool

The JTBD Interview Tool conducts structured customer interviews that reveal the real motivations behind purchasing decisions. Understanding "hiring" and "firing" criteria of products gives product builders the most actionable customer insight possible.

EXERCISE: DESIGN YOUR AI PRODUCT — SUGGESTED RESPONSE

1. **Niche selection example:** "AI-powered proposal writing for freelance consultants." Why: consultants spend 3-5 hours per proposal, proposals directly impact revenue, and the format is structured enough for AI to handle well. Validation: survey 10 consultants about their biggest proposal pain points before building.
2. **JTBD analysis:** Job 1: "When I receive an RFP, I want to quickly understand the client's needs and generate a tailored response, so I can win more contracts without spending my weekend writing." Job 2: "When I finish a project, I want to create a case study from my work, so I can include social proof in future proposals." Job 3: "When a prospect asks for a quote, I want to generate a professional pricing document, so I look established and trustworthy."
3. **MVP brief:** Product: "ProposeAI" — Generate consultant proposals in 10 minutes. Core features: (1) RFP analyzer that extracts requirements, (2) Proposal generator with modular sections, (3) Case study database that auto-inserts relevant proof points. Tech stack: Claude API for generation, FastAPI backend, Next.js frontend, Supabase for storage. Success metric: 10 active users generating 5+ proposals each within 30 days.
4. **Launch plan:** First 10 users: post in freelance consultant communities (r/freelance, Consultant Facebook groups), offer free beta access. Pricing: free for first 3 proposals, then \$29/month. Metrics: proposals generated, time saved per proposal, win rate. Week 1: 5 beta users giving feedback. Week 2: iterate on feedback, fix top 3 issues. Week 4: public launch with 20+ users and first paying customers.
5. **Skill gap analysis:** For this product, key skills needed: Claude API integration (Module 1), FastAPI backend (Module 4), frontend-design skill (Module 5), content generation quality (Module 6). Learning plan: Week 1-2 focus on API integration, Week 3 on UI design, Week 4 on marketing and launch skills.

FINAL ASSESSMENT

20-Question Final Assessment Answers

Below are the answers for all 20 final assessment questions. Each answer includes the correct option and a brief explanation. Use the scoring guide at the bottom to evaluate overall performance.

1. Claude Code's primary interface?

C) Command-line terminal integrated with your repository

Claude Code is a terminal-native tool that works directly within your Git repository.

2. Which agent has "Composer" for multi-file edits?

C) Cursor

Cursor's Composer feature orchestrates multi-file edits from natural language descriptions.

3. LangGraph uses what structure?

C) Directed graphs with nodes and edges

LangGraph represents agent workflows as directed graphs for fine-grained control.

4. CrewAI creates what type of system?

B) Teams of specialized agents with defined roles that collaborate

CrewAI orchestrates multi-agent teams where each agent has a specific role and expertise.

5. What does MCP stand for?

B) Model Context Protocol — standard for connecting AI models to tools

MCP (Model Context Protocol) is Anthropic's open standard for AI-tool integration.

6. Which tool packages LLMs as portable executables?

D) LlamaFile

Mozilla's LlamaFile packages model weights + inference engine into a single executable.

7. vLLM's high-throughput technique?

C) PagedAttention and continuous batching

vLLM uses PagedAttention for KV cache management and continuous batching for GPU efficiency.

8. Highest per-project design rates (\$1,000-10,000)?

C) 3D/Motion design

3D/Motion design commands premium rates due to technical complexity and visual impact.

9. Email marketing ROI per \$1?

C) \$36

Industry data shows *36returnper1* spent on email marketing.

10. GEO optimizes for what type of search?

C) AI-powered answer engines like ChatGPT and Perplexity

GEO (Generative Engine Optimization) targets AI answer engines, not traditional search.

11. Evidence-Based Dialogue ensures what?

C) Claims are verified against reliable sources with confidence levels

This skill requires all claims to be sourced and rated for confidence.

12. JTBD helps discover what about customers?

B) What progress they're trying to make in their lives

JTBD reveals the underlying "job" customers are hiring a product to accomplish.

13. AI Transformation Discovery assessment price?

C) \$5,000-50,000

Enterprise AI assessments command premium consulting fees.

14. Which framework has visual drag-and-drop builder?

C) Dify

Dify provides a visual no-code builder for AI agent workflows.

15. What does Repomix do?

C) Packs entire codebases into a single LLM-optimized file

Repomix creates a single AI-friendly file from an entire repository.

16. 20-hour project in 4 hours at \$100/hr = effective rate?

C) \$500/hour

$2,000 \text{ value} / 4 \text{ hours} = 500/\text{hour}$ effective rate.

17. What makes Unsloth significant?

B) It runs 2-5x faster with 80% less memory, enabling fine-tuning on consumer GPUs

Unsloth's optimizations make model fine-tuning accessible on consumer hardware.

18. Competitive Ads Extractor extracts what insights?

B) Messaging strategies, targeting, offer structures, and creative patterns

The tool analyzes competitor ads for comprehensive strategic intelligence.

19. What is the "knowledge moat" concept?

B) A depth of domain understanding that compounds over time and competitors can't replicate

Knowledge moats are built through repeated, deep research that creates compounding expertise.

20. Common success pattern for AI-native entrepreneurs?

C) Identify pain points, build MVPs in 1-4 weeks with AI, launch fast, iterate

Speed of iteration and learning beats traditional planning-heavy approaches.

Final Assessment Scoring Guide

Score Range	Rating	Interpretation
18-20 correct	GITBUILT MASTER	Exceptional understanding of the AI skills ecosystem. Ready to build and monetize AI-powered products and services immediately.
15-17 correct	ADVANCED BUILDER	Strong command of core concepts. Ready to execute on most AI projects with minimal guidance.
12-14 correct	COMPETENT OPERATOR	Good foundational knowledge. Should review modules where errors occurred and practice the exercises.
8-11 correct	DEVELOPING	Basic awareness but gaps in understanding. Re-read the workbook sections for missed questions and redo exercises.
Below 8	FOUNDATIONAL	Needs significant review. Work through each module sequentially, complete all exercises, and retake the assessment.

INSTRUCTOR TIP: For module quizzes (40 questions total across 8 modules), a score of 32+ (80%) indicates strong module-level comprehension. Students scoring below 24 (60%) on module quizzes should revisit the corresponding modules before attempting the final assessment.

Answers Verified.

Use this answer sheet to assess understanding, guide discussions, and identify areas for deeper exploration. The real test is building something real.