

THE GIT BUILD SERIES

GIT BUILD

WORKBOOK 2

Claude Skills, Workflows & GitHub Repos — The definitive guide to building your AI-native toolkit and turning it into revenue.

GITBUILT DNA + GIT MONEY

THE DNA

THE TOOLS

THE MONEY

Table of Contents

Introduction

Module 1: THE DNA — Coding Agents & IDEs

Module 2: THE DNA — Agent Frameworks

Module 3: THE DNA — MCP Servers & Tooling

Module 4: THE DNA — Local AI & Infrastructure

Module 5: GIT MONEY — Design & Creative Skills

Module 6: GIT MONEY — Marketing & Content Skills

Module 7: GIT MONEY — Research & Intelligence

Module 8: GIT MONEY — Product & Building

Cheat Sheet — All Skills Quick Reference

Final Assessment — 20 Questions

Closing

Introduction

Welcome to **Workbook 2** in the GIT BUILD series. This workbook is the bridge between knowing about AI tools and *actually building* with them. While Workbook 1 laid the foundation of the GIT BUILD philosophy, this workbook goes deeper into the specific skills, frameworks, and repositories that form the DNA of every successful AI-native builder — and more importantly, shows you exactly how each one translates into real income.

The AI landscape in 2025 is defined by three interconnected ecosystems: **coding agents** that write production software, **MCP (Model Context Protocol) servers** that give AI models access to real-world tools and data, and **agent frameworks** that orchestrate complex multi-step workflows. Understanding how these three layers interact is the single most valuable technical skill you can develop right now. Each module in this workbook maps directly to one of these layers, giving you both the conceptual understanding and the hands-on practice to become proficient.

The second half of this workbook — the **GIT MONEY** section — is where things get exciting. We cover the creative, marketing, research, and product skills that turn technical capability into revenue.

Design skills command the highest rates on freelance platforms. Content systems built on AI generate thousands of dollars per month in agency retainers. Research automation gives you competitive intelligence that no manual process can match. And product skills let you ship MVPs in days instead of months. Your GITBUILT DNA is the unique combination of technical depth and business acumen that nobody can copy. Let's build it.

MODULE 1

THE DNA — Coding Agents & IDEs

The 7 foundational tools that write code alongside you

Coding agents represent the most transformative shift in software development since the invention of the compiler. These tools don't just autocomplete your code — they understand your intent, navigate your entire codebase, write multi-file implementations, run tests, debug failures, and iterate on solutions autonomously. The coding agents covered in this module represent the cutting edge of what's possible when you pair large language models with deep code understanding and execution environments. Whether you're a solo founder building a startup MVP, a freelancer delivering projects faster than your competition, or an engineer at a large company looking to 10x your output, these tools are your new team. The key insight is that coding agents are not a replacement for developers — they are a force multiplier. The developers who learn to direct these agents effectively will build more, ship faster, and earn more than those who don't. This module gives you the complete landscape of available tools, their strengths and ideal use cases, and practical exercises to start building with them today.

Claude Code

Anthropic's official command-line coding agent that lives in your terminal and understands your entire repository. Claude Code can read, write, and refactor code across multiple files, run shell commands, execute tests, and iterate on solutions — all through natural language conversation. It supports agentic workflows where it plans, executes, and self-corrects without constant human guidance. This is the most tightly integrated agent for developers who live in the terminal.

<https://docs.anthropic.com/en/docs/claude-code>

AI Coding Agent

Cursor

The AI-native IDE built on VS Code that brings Claude, GPT-4, and other models directly into your editor. Cursor features intelligent codebase-wide search, multi-file editing, inline generation, and a chat interface that understands your project context. Its "Composer" feature lets you describe features in plain English and get complete multi-file implementations. Cursor has become the go-to IDE for AI-first developers, with over 40,000 GitHub stars and rapid adoption across the industry.

<https://cursor.sh> / <https://github.com/getcursor/cursor>

Agent IDE

Codex CLI

OpenAI's open-source command-line agent that brings GPT-4 grade code generation to your terminal. Codex CLI operates as a sandboxed coding assistant that can create files, run commands, and build entire projects from natural language prompts. It supports multiple LLM providers and offers a lightweight, scriptable interface perfect for automation pipelines and CI/CD integration. Codex CLI is ideal for developers who want OpenAI's capabilities without leaving the terminal or switching to a new IDE.

<https://github.com/openai/codex>

AI Coding Agent

Windsurf

Codeium's AI-powered IDE that combines intelligent code completion with agentic workflows. Windsurf features "Cascade," an agent that can understand your codebase context, plan multi-step changes, and execute them across files. It offers real-time collaboration between the developer and the AI, with the agent proposing changes and the developer accepting, modifying, or redirecting as needed. Windsurf is positioned as a direct competitor to Cursor with a focus on speed and a generous free tier.

<https://codeium.com/windsurf>

Agent IDE

Superpowers

A Claude-based coding skill that enhances Claude's native abilities with specialized workflows for common development tasks. Superpowers adds structured prompts, templates, and multi-step agent loops for code review, refactoring, testing, and documentation generation. It acts as a layer on top of Claude that makes the AI more effective at software engineering tasks by providing better context and structured outputs. Superpowers is particularly useful for teams standardizing how they interact with AI coding tools.

<https://github.com/anthropics/claude-code>

AI Coding Skill

Spec Kit

A specification-driven development toolkit that helps you write detailed technical specs before coding, then uses AI to implement them. Spec Kit enforces a "spec-first" workflow where you define what you want to build in structured natural language, and the AI generates implementation plans, code skeletons, and even full implementations. This approach dramatically reduces the iteration cycle and ensures that AI-generated code aligns with your actual requirements. Spec Kit is essential for larger projects where upfront planning prevents expensive rework.

<https://github.com/spec-kit/spec-kit>

Spec-Driven Dev

Aider

An open-source AI pair programming tool that works in your terminal with any Git repository. Aider connects to multiple LLM providers (Claude, GPT-4, local models) and can directly edit files in your repo, creating proper Git commits for each change. It features a chat interface, diff-aware editing, and support for adding images and URLs to the conversation context. Aider is one of the most popular open-source AI coding tools, praised for its clean architecture, transparency, and deep Git integration that makes every AI change traceable.

<https://github.com/paul-gauthier/aider>

AI Pair Programming

\$ GIT MONEY

Coding agents directly translate to income in three powerful ways. First, freelancers using Claude Code or Cursor can deliver projects 3-5x faster, meaning you either earn more per hour or take on more clients. A developer who charges

100/hour but delivers in 4 hours what used to take 20 is effectively earning 500/hour.

Second, startup founders can build and ship MVPs in days instead of months, reducing burn rate and getting to revenue faster. Third, AI-native consulting is a booming market — companies will pay \$150-300/hour for developers who can set up and optimize coding agent workflows for their teams.

EXERCISE

Set Up Your Agent Stack

1. **Choose your primary agent:** Install either Claude Code (`npm install -g @anthropic-ai/claude-code`) or download Cursor from cursor.sh. Spend 15 minutes exploring the interface and reviewing the documentation.
2. **Create a test project:** Initialize a new project directory with a simple structure — a README, a main application file, and a basic configuration file. Use `git init` to set up version control.
3. **Have the agent build something:** Ask your coding agent to build a complete to-do list application with a REST API. Give it a clear prompt that specifies the tech stack (e.g., Python/Flask or Node.js/Express), features (CRUD operations), and any design preferences.
4. **Review and iterate:** Examine the code the agent generated. Ask it to add tests, improve error handling, or add a new feature. Document how many iterations it took and where it struggled.
5. **Document your findings:** Write a brief summary of what worked well, what didn't, and how long the entire process took compared to manual coding. This is your baseline for measuring AI productivity gains.

Write your answer here:

QUIZ

1. Which coding agent is Anthropic's official terminal-based tool that integrates directly with your repository?

- A) Cursor B) Claude Code C) Aider D) Windsurf

2. What is the primary advantage of Cursor's "Composer" feature over traditional code completion?

- A) It only works with Python B) It generates single-line completions C) It creates multi-file implementations from plain English descriptions D) It requires no internet connection

3. Which tool is best described as an open-source AI pair programmer with deep Git integration?

- A) Superpowers B) Spec Kit C) Aider D) Codex CLI

4. What development philosophy does Spec Kit enforce that can reduce iteration cycles?

- A) Test-driven development B) Spec-first workflow C) Behavior-driven development D) Domain-driven design

5. If a freelancer using AI coding agents delivers a project in 4 hours that normally takes 20 hours at \$100/hour, what is their effective hourly rate?

- A) 100 B) 200 C) 400 D) 500

MODULE 2

THE DNA — Agent Frameworks

8 frameworks for building autonomous AI systems

Agent frameworks are the orchestration layer that transforms a single LLM call into a complex, multi-step autonomous workflow. While coding agents help you write software, agent frameworks help you build the software that thinks, decides, and acts on its own. These frameworks provide the plumbing for tool calling, memory management, error recovery, multi-agent collaboration, and human-in-the-loop workflows. The frameworks covered in this module represent the full spectrum of approaches — from lightweight libraries that add just enough structure to raw LLM calls, to full platforms with visual builders, deployment infrastructure, and monitoring dashboards. Choosing the right framework depends on your use case: if you need fine-grained control, LangGraph or pydantic-ai give you maximum flexibility. If you want to ship a production agent quickly, Dify or

CrewAI provide higher-level abstractions. If you're building tools for other developers, CopilotKit or OWL offer specialized interfaces. Understanding this landscape is essential because the agent economy is where the next wave of AI companies will be built. Every SaaS product is becoming an agent product, and these frameworks are the foundation.

OpenClaw

An open-source agent framework designed for building general-purpose AI agents that can use tools, browse the web, and interact with external services. OpenClaw provides a clean architecture for defining agent capabilities, tool integrations, and decision-making loops. It supports multiple LLM backends and offers extensible plugin system for adding new tools and skills. OpenClaw is ideal for developers who want a flexible, community-driven foundation for building custom agents without vendor lock-in.

<https://github.com/openclaw/openclaw>

Agent Framework

LangGraph

Anthropic and LangChain's graph-based framework for building stateful, multi-step AI agents. LangGraph lets you define agent workflows as directed graphs where each node is a function or LLM call, and edges define the control flow. This approach gives you fine-grained control over the agent's decision-making process, including branching logic, loops, parallel execution, and human-in-the-loop checkpoints. LangGraph is the framework of choice for production agents that need reliability, observability, and complex state management.

<https://github.com/langchain-ai/langgraph>

Graph Agent Framework

CrewAI

A multi-agent orchestration framework that lets you define teams of AI agents with specialized roles that collaborate to complete complex tasks. In CrewAI, you create a "crew" of agents — each with their own backstory, expertise, and tools — and define tasks and workflows for how they collaborate. For example, you might have a researcher agent, a writer agent, and an editor agent working together to produce content. CrewAI handles the coordination, delegation, and communication between agents, making it one of the most intuitive frameworks for building collaborative AI systems.

<https://github.com/crewAIInc/crewAI>

Multi-Agent Framework

AutoGPT

The pioneering autonomous AI agent that sparked the agent revolution in 2023. AutoGPT demonstrated that LLMs could be given a high-level goal and autonomously break it down into subtasks, execute them, and self-correct along the way. While the original AutoGPT was experimental, the project has evolved into a mature platform with better memory management, tool use, and reliability. AutoGPT remains one of the most well-known agent frameworks and serves as an excellent learning tool for understanding autonomous agent architecture.

<https://github.com/Significant-Gravitas/AutoGPT>

Autonomous Agent

Dify

An open-source LLM application development platform that provides a visual workflow builder for creating AI agents, RAG pipelines, and AI-powered applications. Dify features a drag-and-drop interface for designing agent workflows, built-in RAG capabilities with document ingestion, prompt management, model switching, and deployment tools. It supports over 50 LLM providers and offers both cloud and self-hosted options. Dify is the fastest path from idea to deployed AI application, especially for teams that include non-developers in the workflow design process.

<https://github.com/langgenius/dify>

Visual Agent Platform

OWL

A lightweight, opinionated agent framework designed for building focused, task-specific agents with minimal boilerplate. OWL emphasizes simplicity and developer experience, providing just enough structure to build reliable agents without the complexity of larger frameworks. It includes built-in support for tool calling, structured output, and conversation management. OWL is ideal for developers who find larger frameworks too heavy and want to build agents that do one thing extremely well rather than trying to be general-purpose.

<https://github.com/owl-framework/owl>

Lightweight Agent

CopilotKit

A framework for building AI copilot experiences directly into web applications. CopilotKit provides React components, hooks, and a runtime that make it straightforward to add chat-based AI assistants, autonomous agents, and AI-powered features to any web app. It handles the complex parts of copilot development — context management, tool integration, streaming responses, and state synchronization — so you can focus on building the user experience. CopilotKit is the go-to choice for SaaS companies looking to add AI capabilities to their existing products.

<https://github.com/CopilotKit/CopilotKit>

Copilot Framework

pydantic-ai

A Python agent framework built by the team behind Pydantic, leveraging type-safe models for structured agent interactions. pydantic-ai uses Pydantic models to define agent inputs, outputs, and tool interfaces, providing automatic validation, serialization, and IDE autocomplete for everything your agent does. This type-safe approach catches errors at development time rather than runtime, making agent development more reliable and maintainable. pydantic-ai is the natural choice for Python developers who value type safety and clean architecture in their agent code.

<https://github.com/pydantic/pydantic-ai>

Type-Safe Agent

\$ GIT MONEY

Agent-powered SaaS is the single biggest AI business opportunity in 2025. The pattern is clear: take a manual, time-consuming business process, build an agent that automates 80% of it, and sell access. Customer support agents, research assistants, content pipelines, data processing workflows — every industry has processes begging for automation. Using frameworks like CrewAI or Dify, a solo developer can build and ship an agent-powered product in 2-4 weeks. The key insight is that customers don't care about your framework choice — they care about the result. Price based on value delivered, not hours worked. An agent that saves a company 50,000/year in manual labor can easily command 500-1,000/month in subscription revenue.

EXERCISE

Design Your First AI Agent

1. **Identify a real problem:** Think of a repetitive, time-consuming task you or someone you know does regularly. Write down the task, who does it, how long it takes, and what the output looks like. The more specific, the better your agent will be.
2. **Choose your framework:** Based on the complexity of the task, select a framework. For a single-task agent, use pydantic-ai or OWL. For a multi-step workflow, use LangGraph. For a collaborative multi-agent system, use CrewAI. For rapid prototyping with visual design, use Dify. Write down your choice and one sentence explaining why.
3. **Outline the architecture:** Draw or describe the agent's decision flow. What inputs does it receive? What tools does it need (web search, API calls, file operations)? What are the decision points? What is the final output format? Include at least three steps in your workflow.
4. **Define the tool interfaces:** For each tool your agent will use, write a simple specification: tool name, description, input parameters, and expected output. Use Pydantic model syntax if you're comfortable with it, or plain English otherwise.
5. **Plan the MVP:** List the minimum set of features needed to demonstrate your agent works. Estimate how long each would take to build with AI assistance. Set a target date for a working prototype.

Write your answer here:

QUIZ

1. Which framework uses directed graphs to define agent workflows with fine-grained control over state and decision-making?

- A) CrewAI B) LangGraph C) AutoGPT D) Dify

2. What is the primary value proposition of CrewAI's multi-agent approach?

- A) It uses the smallest models possible B) It enables specialized agents with defined roles to collaborate on complex tasks C) It only works with Claude D) It replaces the need for any human input

3. Which framework would be best for a non-technical team member to design an AI workflow visually?

- A) pydantic-ai B) LangGraph C) Dify D) OWL

4. What makes pydantic-ai unique among agent frameworks?

- A) It uses Pydantic models for type-safe agent interactions B) It only works with vision models
C) It was the first agent framework ever created D) It requires no Python knowledge

5. An agent that saves a company \$50,000/year could reasonably be priced at what monthly subscription rate?

- A) 10/month B) 50/month C) 500 – 1,000/month D) 5,000/month

MODULE 3

THE DNA — MCP Servers & Tooling

7 tools extending AI capabilities through the Model Context Protocol

The Model Context Protocol (MCP) is rapidly becoming the standard way to connect AI models to external tools, data sources, and services. Think of MCP as the "USB-C of AI" — a universal connector that lets any AI model plug into any tool or data source through a standardized interface. Before MCP, every AI integration was a custom build: you'd write specific code to connect Claude to your database, GPT to your API, and so on. MCP eliminates this redundancy by providing a protocol that any model can use to discover and interact with any tool. The MCP servers covered in this module represent the most useful and well-built tools in the ecosystem. Some provide access to external services (Tavily for web search, Context7 for documentation), some provide capabilities

(Playwright for browser automation, markdownify for content conversion), and some provide infrastructure (fastmcp for building your own servers, MCPHub for discovery). Understanding MCP is essential because it's the infrastructure layer that makes AI agents truly useful — an agent without tools is just a chatbot, but an agent with MCP servers can search the web, read documentation, automate browsers, and interact with any API.

Tavily MCP Server

A Model Context Protocol server that provides AI models with real-time web search capabilities through the Tavily search API. Tavily is specifically designed for AI agents rather than human search, returning clean, structured results optimized for LLM consumption rather than SEO-optimized web pages. The MCP server exposes search, extract, and crawl tools that agents can use to find current information, extract content from specific URLs, and discover related resources. Tavily is the de facto standard for adding web search to AI agent workflows.

<https://github.com/tavily-ai/tavily-mcp>

MCP — Web Search

Context7

An MCP server that provides AI models with up-to-date library and framework documentation on demand. Instead of relying on training data that may be outdated, Context7 fetches the latest documentation for any supported library when the agent needs it. This means your agent can reference current APIs, syntax, and best practices for any library without you needing to manually paste documentation into the context. Context7 supports hundreds of libraries and frameworks, making it an essential tool for coding agents that work across multiple technology stacks.

<https://github.com/upstash/context7-mcp>

MCP — Documentation

Task Master AI

An MCP server and task management system that helps AI agents break down complex projects into manageable tasks, track progress, and maintain context across long development sessions. Task Master AI provides tools for creating, updating, and querying tasks, as well as generating implementation plans based on project requirements. It's particularly valuable for large projects where maintaining context across multiple coding sessions is critical. The server integrates with Claude Code and other MCP-compatible agents to provide persistent task management.

<https://github.com/eyaltoledano/claude-task-master>

MCP — Task Management

MCP Playwright

An MCP server that gives AI agents the ability to control a web browser through the Playwright automation library. With MCP Playwright, agents can navigate to websites, click elements, fill forms, take screenshots, extract content, and interact with web applications just like a human user would. This is essential for tasks like web scraping, automated testing, form filling, and any workflow that requires browser interaction. The server provides tools for navigation, element selection, page interaction, and screenshot capture.

<https://github.com/anthropics/mcp-playwright>

MCP — Browser Automation

fastmcp

A Python library for rapidly building custom MCP servers with minimal boilerplate. fastmcp provides decorators and utilities that let you turn any Python function into an MCP tool with a single line of code. It handles the protocol details, serialization, and server infrastructure so you can focus on defining what your tools do. fastmcp is the fastest way to create MCP servers, making it the go-to choice for developers who want to expose their own tools, APIs, or data sources to AI agents. It supports both stdio and SSE transport.

<https://github.com/jlowin/fastmcp>

MCP — Server Builder

markdownify-mcp

An MCP server that converts web content from any URL into clean, structured Markdown format. This is crucial for AI agents that need to process web content but struggle with raw HTML. markdownify-mcp fetches web pages, strips away navigation, ads, and other noise, and returns the core content as well-formatted Markdown. This makes web content immediately usable by LLMs without additional preprocessing. The server handles JavaScript-rendered pages, respects robots.txt, and provides configurable content extraction options.

<https://github.com/markdownify-mcp/markdownify-mcp>

MCP — Content Conversion

MCPHub

A centralized discovery platform and registry for MCP servers, making it easy to find, evaluate, and install MCP servers for your AI workflow. MCPHub provides a searchable catalog of community-built and official MCP servers with ratings, documentation links, and installation instructions. Think of it as the "npm for MCP servers" — a single place to discover tools that extend your AI agent's capabilities. MCPHub also provides tools for managing multiple MCP server configurations and keeping your servers up to date.

<https://github.com/mcphub/mcphub>

MCP — Discovery Registry

\$ GIT MONEY

Building and selling MCP servers is an emerging business model with enormous potential. Every company that adopts AI agents will need custom MCP servers to connect those agents to their internal tools, databases, and workflows. A single well-built MCP server that solves a common integration problem can be sold as a SaaS product or licensed to enterprises. Consider the math: if you build an MCP server that connects AI agents to Salesforce, and charge *49/month per user, you need only 200 customers to hit 10K MRR*. The MCP ecosystem is still early — first movers who build high-quality servers for popular platforms will establish lasting advantages.

EXERCISE

Build a Custom MCP Server

- 1. **Identify a tool gap:** Think of a tool or data source that AI agents currently can't access but would find valuable. This could be an internal API, a niche data source, a specific file format parser, or a service without an existing MCP server. Write down the problem and why existing solutions don't cover it.
- 2. **Define your server's tools:** List 3-5 specific tools your MCP server will expose. For each tool, write: the tool name, a clear description of what it does, the input parameters (with types), and the expected output format. Use a table format for clarity.
- 3. **Choose your build approach:** Decide whether to use fastmcp (Python) for rapid development or build from scratch using the MCP SDK. Write a brief justification for your choice, considering your familiarity with the language and the complexity of the tools.
- 4. **Write the specification:** Create a formal specification document for your MCP server that includes: server name and description, list of tools with full schemas, any resources or prompts the server provides, and the transport method (stdio or SSE). This spec should be detailed enough that another developer could build it.
- 5. **Plan your go-to-market:** How would you distribute this MCP server? Options include: open source with paid support, freemium with premium features, enterprise licensing, or listing on MCPHub. Choose one approach and write a brief launch plan.

Write your answer here:

QUIZ

1. What analogy best describes the Model Context Protocol's role in the AI ecosystem?

- A) A programming language B) A universal connector (like USB-C) for AI models and tools C) A specific LLM model D) A cloud hosting platform

2. Why is Tavily's search API particularly suited for AI agents compared to traditional search engines?

- A) It's faster than Google B) It returns results optimized for LLM consumption rather than human browsing C) It only searches social media D) It requires no API key

3. What problem does Context7 solve for coding agents?

- A) It speeds up code compilation B) It provides up-to-date library documentation on demand instead of relying on outdated training data C) It writes unit tests automatically D) It deploys code to production

4. Which library is specifically designed for rapidly building custom MCP servers with minimal code?

- A) Playwright B) markdownify C) fastmcp D) Tavily

5. If you built an MCP server charging $49/\text{month}$ and wanted to reach 10K MRR, how many customers would you need?

- A) 50 B) 100 C) 204 D) 500

MODULE 4

THE DNA — Local AI & Infrastructure

7 tools for running AI locally and at scale

Running AI models locally is one of the most strategically important capabilities you can develop. While cloud-based APIs from Anthropic, OpenAI, and Google are powerful and convenient, they come with recurring costs, latency from network round-trips, data privacy concerns, and dependency on external service availability. Local AI infrastructure eliminates all of these issues. The tools in this module span the full spectrum from running a single model on your laptop to serving models at production scale. Ollama and LlamaFile make it trivial to run models locally — literally one command to download and start a model. Open WebUI gives you a ChatGPT-like interface for your

local models. Unsloth lets you fine-tune models efficiently on consumer hardware. vLLM and FastAPI provide production-grade serving infrastructure. Portkey Gateway adds the observability, caching, and load balancing layer that production systems need. The strategic advantage of local AI is compound: every model you run locally is a model you don't pay per-token costs for. Over months and years of usage, this adds up to enormous savings. For startups and indie developers, this can be the difference between profitability and burning through runway on API costs alone.

Ollama

The most popular tool for running large language models locally with a single command. Ollama provides a streamlined experience for downloading, running, and managing LLMs on your own hardware — Mac, Linux, or Windows. With simple commands like "ollama run llama3" or "ollama run codellama," you can have a capable AI model running on your machine in minutes. Ollama supports hundreds of models from the GGUF ecosystem, handles quantization automatically, and provides an API compatible with OpenAI's format, making it a drop-in replacement for cloud APIs in most applications.

<https://github.com/ollama/ollama>

Local LLM Runtime

Open WebUI

A full-featured, self-hosted web interface for AI models that provides a ChatGPT-like experience for your local Ollama models or remote API endpoints. Open WebUI supports chat, file uploads, image generation, web search integration, user management, and conversation history — all running on your own hardware with full data privacy. It features a beautiful, responsive UI, support for multiple models, RAG capabilities with document upload, and an extensible plugin system. Open WebUI is the go-to choice for anyone who wants a production-quality AI chat interface without sending data to third parties.

<https://github.com/open-webui/open-webui>

Local AI Interface

LlamaFile

Mozilla's innovation in AI distribution — a tool that packages an LLM into a single executable file that runs on any computer without installation. LlamaFile embeds the model weights and inference engine into one binary that you can share via email, USB drive, or download link. The recipient simply runs the file and gets a local AI chat server in their browser. This approach eliminates all the complexity of model installation, Python environments, and dependency management. LlamaFile is revolutionary for distributing AI capabilities to non-technical users and for air-gapped or offline environments.

<https://github.com/Mozilla-Ocho/llamafile>

Portable AI

Unsloth

A library for fine-tuning large language models 2-5x faster with 80% less memory usage compared to standard Hugging Face training. Unsloth achieves this through custom Triton kernels, optimized attention mechanisms, and smart memory management that makes it possible to fine-tune models on consumer GPUs — even on a single RTX 3090 or Apple M-series chip. Unsloth supports LoRA and QLoRA fine-tuning for models like Llama, Mistral, and Gemma. It's the essential tool for anyone who wants to create custom, domain-specific models without enterprise-grade hardware budgets.

<https://github.com/unslothai/unsloth>

Model Fine-Tuning

vLLM + FastAPI

A production-grade model serving stack that combines vLLM's high-throughput inference engine with FastAPI's lightweight web framework. vLLM uses PagedAttention and continuous batching to serve models 5-10x faster than standard Hugging Face inference, while using significantly less memory. Combined with FastAPI, you get a REST API endpoint for your model that handles concurrent requests, streaming responses, and health checks. This stack is the standard for deploying local models to production, whether serving internal teams or external customers.

<https://github.com/vllm-project/vllm>

Model Serving

Portkey Gateway

An AI gateway that provides observability, caching, load balancing, and fallback management for AI API calls. Portkey sits between your application and AI providers (Anthropic, OpenAI, local models), giving you unified logging, latency tracking, cost monitoring, and automatic failover between providers. It supports caching responses to reduce costs and latency, rate limiting, and A/B testing between models. Portkey Gateway is the infrastructure layer that makes AI applications reliable and cost-efficient in production, especially when using multiple model providers or running hybrid cloud/local setups.

<https://github.com/portkey-ai/gateway>

AI Gateway

\$ GIT MONEY

Running AI locally is not just a technical choice — it's a financial strategy. Consider a content generation SaaS that processes 1 million tokens per day. At Claude's API pricing, that's roughly 15 / day or 450 / month in API costs alone. Running the same workload on a local model with Ollama costs 0 in token fees — you only pay for the hardware, which is a one-time investment. Over a year, that's \$5,400+ in savings that goes directly to your bottom line. For startups operating on tight margins, local AI can be the difference between profitability and running out of runway. The highest-margin AI businesses are the ones that minimize their API dependency.

EXERCISE

Deploy a Local AI Stack

1. **Install Ollama:** Download and install Ollama from ollama.ai for your operating system. Verify the installation by running "ollama list" in your terminal. If you encounter any issues, document them and how you resolved them.
2. **Download and run a model:** Use "ollama pull llama3.1" to download Meta's Llama 3.1 model (or another model of your choice). Then run "ollama run llama3.1" to start an interactive chat session. Ask the model three questions and note the quality of responses compared to Claude or GPT-4.
3. **Set up Open WebUI:** Install and configure Open WebUI following the documentation at their GitHub repository. Connect it to your local Ollama instance. Navigate the interface and test at least three features: chat, file upload, and model selection.
4. **Connect an application:** Write a simple Python script or use an existing tool (like Claude Code or Aider) that connects to your local model through Ollama's OpenAI-compatible API endpoint. Demonstrate that it can generate code, answer questions, or perform another useful task.
5. **Calculate your savings:** Estimate your typical monthly AI API usage in tokens. Using Ollama's local inference, calculate how much you would save per month and per year. Document your calculation and the hardware you're using.

Write your answer here:

QUIZ

1. What command would you use to download and start running Llama 3.1 with Ollama?

- A) ollama install llama3.1 B) ollama pull llama3.1 then ollama run llama3.1 C) pip install ollama-llama3.1 D) ollama download llama3.1

2. What is LlamaFile's key innovation in AI model distribution?

- A) It makes models smaller B) It packages a model into a single executable file that runs without installation C) It converts models to JavaScript D) It trains models faster

3. How does Unsloth achieve faster fine-tuning with less memory?

- A) By using smaller models only B) Through custom Triton kernels, optimized attention, and smart memory management C) By running in the cloud D) By using fewer training examples

4. What technique does vLLM use to achieve 5-10x faster inference than standard serving?

- A) Model compression B) PagedAttention and continuous batching C) Distributed training D) Edge computing

5. A content SaaS processing 1M tokens/day at ~\$15/day in API costs would save approximately how much per year by switching to local inference?

- A) 1,200 B) 3,600 C) 5,400 D) 15,000

MODULE 5

GIT MONEY — Design & Creative Skills

8 skills for AI-powered design that commands premium rates

Design is the highest-demand AI skill category in 2025, and it's not even close. The reason is simple: every business needs design, most businesses can't afford dedicated designers, and AI can now produce design work that's 80-90% as good as a mid-level professional. That last 10-20% gap is where skilled AI operators create their competitive advantage. The design skills in this module cover the full spectrum from frontend UI design to 3D rendering, from algorithmic art to hand-drawn diagrams. Each skill gives you the ability to produce a specific type of design output that businesses will pay for. Frontend design is the bread and butter — landing pages, app interfaces, email templates. Canvas design covers print-ready materials, social media graphics, and

marketing collateral. Algorithmic art and 3D/motion skills command the highest per-project rates because they require the most creative direction. The key insight for monetization is that AI doesn't replace designers — it replaces design agencies. A single skilled operator with these AI tools can produce the output of an entire design team. This means you can offer agency-quality work at freelancer prices and still maintain excellent margins. The freelancers and agencies who adopt these tools first will capture market share from those who don't.

frontend-design \$

A Claude skill that generates production-ready frontend UI designs from text descriptions. Specify the type of page (landing page, dashboard, e-commerce, portfolio), the style (minimal, bold, corporate, playful), and the content, and the skill produces complete HTML/CSS with modern design principles. Outputs include responsive layouts, proper typography hierarchy, color systems, and component-based structures. This skill is the foundation of an AI-powered web design service — you can generate complete website designs in minutes and charge \$500-2,000 per project.

<https://github.com/anthropics/claude-skills>

UI/UX Design

canvas-design \$

A design skill that creates visual compositions optimized for print and digital canvas formats. Canvas-design generates layouts for presentations, social media posts, infographics, business cards, and marketing materials. It understands design principles like the rule of thirds, visual hierarchy, whitespace management, and brand consistency. The skill outputs designs that can be directly used or further refined in tools like Figma or Canva. This is your tool for creating client-ready marketing materials without a graphic design degree.

<https://github.com/anthropics/claude-skills>

Visual Design

algorithmic-art \$\$

A generative art skill that creates unique visual artworks using mathematical algorithms, code-driven patterns, and procedural generation techniques. Algorithmic-art produces pieces that range from geometric abstractions and fractal patterns to data visualizations and organic forms. Each output is unique and can be generated at various resolutions suitable for digital display, print, or NFT minting. This skill commands premium rates because algorithmic art sits at the intersection of technology and fine art — collectors and brands pay \$200-5,000 for commissioned generative pieces.

<https://github.com/anthropics/claude-skills>

Generative Art

Color Expert \$

An AI skill specialized in color theory, palette generation, and brand color system design. Color Expert generates harmonious color palettes based on psychology, industry conventions, and accessibility standards. It can analyze existing brand colors and suggest complementary, analogous, or triadic extensions. The skill ensures all palettes meet WCAG contrast requirements for accessibility compliance — a critical requirement for enterprise clients. Use this as an add-on service for any design project or as a standalone brand identity consultation offering.

<https://github.com/anthropics/claude-skills>

Color & Branding

Hand-Drawn Diagrams \$

A skill that creates diagrams with a hand-drawn, sketch-like aesthetic that feels approachable and human. Unlike sterile, perfectly geometric diagrams, hand-drawn diagrams use slightly irregular lines, handwritten-style text, and an informal visual language that's perfect for presentations, documentation, educational materials, and pitch decks. This style is particularly popular in tech startups and educational content because it communicates ideas clearly without feeling overly corporate. It's a unique differentiator that sets your deliverables apart from template-based design tools.

<https://github.com/anthropics/claude-skills>

Diagram Design

claude designs skills (3D/Motion) \$\$\$

A comprehensive design skill collection that covers 3D rendering, motion graphics, and advanced visual effects. This skill set generates 3D scenes, product visualizations, animated elements, and motion design specifications that can be rendered in tools like Blender, Three.js, or After Effects. The 3D/motion skill category commands the highest rates in AI design because it requires the most creative direction and produces the most visually impressive outputs. Product renders, animated logos, and 3D social media content can command \$1,000-10,000 per project from brands and agencies.

<https://github.com/anthropics/claude-skills>

3D / Motion

Nothing Design \$

A minimalist design philosophy and skill that focuses on the art of reduction — creating powerful visual communication through simplicity, negative space, and restraint. Nothing Design generates clean, impactful layouts that follow the "less is more" principle, stripping away decorative elements to focus on essential messaging. This approach is highly valued in luxury branding, tech company landing pages, and editorial design. Clients who want their brand to feel premium, sophisticated, and intentional specifically seek out this minimalist aesthetic.

<https://github.com/anthropics/claude-skills>

Minimalist Design

GPT Image 2 \$\$

OpenAI's advanced image generation model that produces photorealistic images, illustrations, and design elements from text descriptions. GPT Image 2 represents the state of the art in AI image generation, capable of producing images with consistent style, accurate text rendering, and complex compositions. Combined with Claude's design skills for layout and structure, GPT Image 2 fills the need for custom photography, product mockups, illustration, and visual content that would otherwise require a professional photographer or illustrator. This combination creates a complete AI design studio.

<https://openai.com/index/gpt-image>

Image Generation

DNA STRAND

The creative DNA — design skills are the highest-demand AI skill category. Every business needs visual content, and AI can produce it at unprecedented speed and scale. The operators who master these skills don't just save time — they unlock entirely new business models. A single person with these tools can run a design agency, create and sell templates, produce social media content for dozens of clients simultaneously, and generate passive income through digital product sales. This is the DNA that turns technical AI knowledge into creative output that people will pay for.

EXERCISE

Build a Design Portfolio in 30 Minutes

1. **Set up your workspace:** Open Claude or your preferred AI tool with the frontend-design skill loaded. Create a new folder called "ai-design-portfolio" to organize your outputs.
2. **Generate Design #1 — SaaS Landing Page:** Use the frontend-design skill to create a complete landing page for a fictional SaaS product. Specify: product name, target audience, key features (3-4), pricing tier, and visual style (choose one: minimal, bold, or corporate). Save the output as an HTML file.
3. **Generate Design #2 — Portfolio Website:** Create a personal portfolio website for a fictional creative professional (photographer, designer, or writer). Include: hero section, work showcase (4-6 items), about section, and contact form. Choose a distinctly different visual style from Design #1.
4. **Generate Design #3 — E-commerce Product Page:** Design a product detail page for an e-commerce store selling a physical product (choose any product). Include: product image area, description, specifications, reviews section, and related products. Use a third visual style.
5. **Critique and compare:** Review all three designs side by side. For each, write: what worked well, what could be improved, and how long it took to generate. Rate each design on a 1-10 scale for visual appeal, usability, and commercial viability. This comparison will help you understand the strengths and limitations of AI-generated design.

Write your answer here:

QUIZ

1. Why is design considered the highest-demand AI skill category in 2025?

- A) It's the easiest to learn B) Every business needs design, most can't afford designers, and AI can produce 80-90% quality output C) Only AI can do design D) Design tools are free

2. What makes the Hand-Drawn Diagrams skill uniquely valuable for certain audiences?

- A) It's faster than other design tools B) Its sketch-like, approachable aesthetic feels human and is popular in startups and education C) It uses no colors D) It only works for technical diagrams

3. Which design skill category commands the highest per-project rates and why?

- A) Frontend design — most demand B) Color Expert — rarest skill C) 3D/Motion — requires most creative direction and produces most impressive outputs D) Canvas design — most versatile

4. What is the strategic insight about AI and the design agency model?

- A) AI will eliminate all design agencies B) A single skilled operator with AI tools can produce agency-quality output, replacing the need for large teams C) Agencies will only use AI for minor tasks D) AI design is always inferior to human design

5. What accessibility standard does the Color Expert skill ensure compliance with?

- A) GDPR B) SOC 2 C) WCAG contrast requirements D) HIPAA

MODULE 6

GIT MONEY — Marketing & Content Skills

8 skills powering the 350K-follower content machine

Content is the engine of online business, and AI has completely transformed how content is created, optimized, and distributed. The skills in this module represent a complete content marketing system — from understanding algorithms and generating hooks that stop the scroll, to creating full social media posts, optimizing them for maximum engagement, and extracting competitive intelligence from paid advertising. This isn't about spamming generic AI content. This is about using AI as a force multiplier for a strategic content operation. The creator behind these skills built a following of

over 350,000 people across platforms using this exact system. The key is that these skills work together as a pipeline: the voice-builder establishes your brand voice, the hook-generator creates attention-grabbing opening lines, the post-scorer predicts engagement before you publish, and the algorithm optimizer ensures your content gets maximum distribution. Each skill in isolation is useful, but the real power is in the system. For monetization, content skills open doors to freelance content creation (2,000 – 10,000/*monthperclient*), *contentagency* services (10,000-50,000/month), personal brand building (which compounds into consulting, courses, and product sales), and competitive intelligence services for businesses. The content DNA you build here is the skill set that keeps generating returns long after you've mastered it.

Social Media OS (Charlie Hills) \$\$\$

A comprehensive operating system for social media management that integrates content planning, creation, scheduling, and analytics into a single AI-powered workflow. Social Media OS helps you maintain a consistent posting schedule across multiple platforms, adapt content for each platform's unique format and audience, and track performance metrics to continuously improve your strategy. The system includes content calendars, audience research tools, and performance prediction models. This is the backbone skill for anyone running a content operation — personal brand, agency, or business social media presence.

<https://github.com/charliehills/social-media-os>

Content System

hook-generator \$\$

An AI skill specifically designed to create attention-grabbing hooks — the opening lines or frames that determine whether someone stops scrolling or keeps going. The hook-generator understands the psychology of attention across platforms (Twitter/X, LinkedIn, Instagram, TikTok, YouTube) and generates hooks optimized for each. It uses proven frameworks like curiosity gaps, contrarian statements, specific numbers, story openings, and pattern interrupts. A great hook can increase content engagement by 300-500%, making this skill the single highest-ROI content optimization tool in the stack.

<https://github.com/anthropics/claude-skills>

Content Creation

post-scorer \$

An AI-powered content scoring system that predicts how well a social media post will perform before you publish it. The post-scorer analyzes your draft content against engagement patterns, platform algorithm preferences, audience demographics, and timing factors to generate a predicted engagement score. It provides specific suggestions for improvement: word choice, structure, length, call-to-action, and formatting. Using post-scorer before publishing can significantly improve your content's reach and engagement, making every piece of content more valuable.

<https://github.com/anthropics/claude-skills>

Content Optimization

youtube-thumbnail \$\$

A specialized design skill for creating YouTube thumbnails that maximize click-through rates. The youtube-thumbnail skill understands what makes thumbnails effective on YouTube's platform — high contrast, expressive faces, bold text overlays, curiosity elements, and brand consistency. It generates thumbnail concepts, suggests compositions, and can produce complete thumbnail designs optimized for both desktop and mobile viewing. Since thumbnails are one of the top two factors in YouTube video performance (along with title), this skill directly impacts video views and channel growth.

<https://github.com/anthropics/claude-skills>

YouTube Growth

Email Marketing Bible \$\$\$

A comprehensive email marketing skill that covers the entire email lifecycle — from subject line generation and copywriting to segmentation, automation, and analytics. The Email Marketing Bible includes frameworks for welcome sequences, promotional campaigns, re-engagement flows, and newsletter strategies. It understands email deliverability best practices, spam filter avoidance, and the psychology of email engagement. Email marketing remains the highest-ROI digital marketing channel, with an average return of 36 *forevery*1 spent. This skill helps you capture that value.

<https://github.com/anthropics/claude-skills>

Email Marketing

GEO/SEO Claude \$\$

An AI skill for Generative Engine Optimization (GEO) and Search Engine Optimization (SEO) that helps content rank in both traditional search results and AI-powered answer engines. As more people use ChatGPT, Perplexity, and Google's AI overviews to find information, traditional SEO is no longer enough. GEO/SEO Claude optimizes content for AI citation, structured data, entity recognition, and the specific factors that AI search engines use to select and rank content. This dual-optimization approach ensures your content is discoverable across all search paradigms.

<https://github.com/anthropics/claude-skills>

SEO / GEO

Twitter Algorithm Optimizer **\$\$**

A skill that analyzes and optimizes content specifically for the Twitter/X algorithm. The optimizer understands the platform's ranking factors — engagement velocity, reply quality, retweet patterns, bookmark signals, and follower interaction — and provides specific recommendations to maximize reach. It can analyze your existing tweets to identify patterns in your high-performing content, suggest optimal posting times based on your audience, and generate tweet variants optimized for algorithmic amplification. For anyone building a personal brand or business on Twitter/X, this skill is a competitive advantage.

<https://github.com/anthropics/claude-skills>

Twitter/X Growth

Competitive Ads Extractor **\$\$\$**

An intelligence-gathering skill that analyzes competitor advertising campaigns to extract strategic insights. The Competitive Ads Extractor can process ad libraries, landing pages, and ad creative to identify competitor messaging strategies, targeting approaches, offer structures, and creative patterns. This intelligence is invaluable for businesses planning their own advertising — you can learn from competitors' tested campaigns rather than starting from scratch. Agencies charge \$2,000-10,000 for competitive ad analysis reports that this skill can produce in hours.

<https://github.com/anthropics/claude-skills>

Competitive Intel

DNA STRAND

Content DNA — the 17-skill system behind 350K followers. The power of these marketing and content skills isn't in any individual tool — it's in the pipeline they create together. Voice builder establishes your brand identity. Hook generator captures attention. Post scorer predicts performance. Algorithm optimizer maximizes distribution. Email Marketing Bible captures and nurtures leads. GEO/SEO Claude ensures discoverability. Together, these skills form a content machine that compounds over time, growing your audience and revenue with each piece of content you publish.

EXERCISE

Create a Content Machine

1. **Define your content persona:** Choose a niche and create a content persona — a fictional representation of your target audience. Include: demographics, pain points, aspirations, where they spend time online, and what type of content they engage with. Write at least 200 words defining this persona.
2. **Generate three hooks:** Using the hook-generator skill or your own AI tool, create three different hooks for the same core idea. Use three different frameworks: one curiosity gap hook, one contrarian statement hook, and one specific numbers/data hook. Write all three out.
3. **Score and iterate:** Take your best hook and expand it into a full social media post (200–300 words). Then use the post-scorer approach — evaluate the post on: attention capture, value delivery, readability, call-to-action strength, and shareability. Rewrite the post incorporating your improvements.
4. **Plan a week of content:** Using the Social Media OS framework, plan 5 posts for the week. For each post, write: the platform, the hook, the core message, the format (text, image, video, thread), and the call-to-action. Ensure variety in format and topic.
5. **Measure and reflect:** Write a brief analysis of what you learned from this exercise. What was hardest? What surprised you? How would you price this content creation service if you were offering it to a client?

Write your answer here:

QUIZ

1. What is the estimated engagement increase from using an optimized hook compared to a generic opening?

- A) 10-20% B) 50-100% C) 300-500% D) 1000%+

2. What does GEO stand for and why is it important beyond traditional SEO?

- A) Google Engine Optimization — it helps with Google ads B) Generative Engine Optimization — it optimizes for AI-powered answer engines like ChatGPT and Perplexity C) General Email Outreach — it improves email marketing D) Global Engagement Optimization — it increases social media reach

3. What is the average ROI for email marketing according to industry data?

- A) 5 *forevery*1 spent B) 12 *forevery*1 spent C) 36 *forevery*1 spent D) 100 *forevery*1 spent

4. Which two factors are the top determinants of YouTube video performance?

- A) Video length and production quality B) Upload time and subscriber count C) Thumbnail and title D) Number of tags and description length

5. How much do agencies typically charge for competitive ad analysis reports?

- A) 100 – 500 B) 500-1,000 C) 2,000 – 10,000 D) 50,000+

MODULE 7

GIT MONEY — Research & Intelligence

5 skills for deep research that compounds into competitive advantage

Research is the silent killer app of AI. While everyone is focused on content creation and code generation, the most financially valuable AI application for many businesses is deep research — the ability to rapidly synthesize information from dozens or hundreds of sources into actionable intelligence. The research skills in this module represent different facets of AI-powered research: from comprehensive web research engines that can process vast amounts of information, to academic research tools that navigate scholarly literature, to evidence-based dialogue systems that ensure conclusions are grounded in facts rather than speculation. The competitive advantage of research skills is that they compound. Every research project you complete makes you more

knowledgeable, which makes your next research project faster and more valuable. Over time, this creates a knowledge moat that competitors cannot easily replicate. A consultant who can deliver a comprehensive market analysis in 2 hours using AI research tools, when it used to take 2 weeks manually, has an enormous competitive advantage. They can serve more clients, charge premium rates for faster turnaround, and provide deeper analysis because the AI handles information gathering while they focus on strategic interpretation. The research skills here are your foundation for building that advantage.

Deep Research Engine \$\$\$

A multi-step AI research system that goes far beyond a single search query. The Deep Research Engine takes a research question, decomposes it into sub-questions, searches across multiple sources, evaluates source credibility, synthesizes findings, identifies gaps, and produces a comprehensive research report with citations. It iteratively refines its search based on what it finds, following leads and exploring related topics to ensure thoroughness. The output is a structured report with executive summary, detailed findings, source list, and confidence assessments. This is the tool for producing research that would normally require a team of analysts.

<https://github.com/anthropics/claude-skills>

Deep Research

Academic Research Skills \$\$\$

A specialized research skill for navigating and synthesizing academic literature. Academic Research Skills can search across databases like arXiv, PubMed, Google Scholar, and Semantic Scholar to find relevant papers, extract key findings, identify research trends, and produce literature reviews. It understands academic citation formats, research methodology, and the structure of scientific papers. This skill is invaluable for graduate students, researchers, and consultants who need to stay current with scientific literature or produce evidence-based reports. It turns weeks of literature review into hours of AI-assisted synthesis.

<https://github.com/anthropics/claude-skills>

Academic Research

Evidence-Based Dialogue \$\$

A research skill that ensures AI-driven conversations and conclusions are grounded in verifiable evidence rather than speculation or common misconceptions. Evidence-Based Dialogue cross-references claims against reliable sources, flags unsupported assertions, provides confidence levels for different statements, and suggests areas where more research is needed. This skill is particularly valuable for consultants, analysts, and decision-makers who need to ensure their AI-assisted research meets professional standards of accuracy and reliability. It acts as a quality control layer for all AI-generated research outputs.

<https://github.com/anthropics/claude-skills>

Research Quality

Social Media Research \$\$

A research skill focused on extracting insights from social media platforms — analyzing trends, sentiment, influencer networks, and community discussions. Social Media Research can monitor brand mentions, track conversation themes, identify emerging trends before they go mainstream, and analyze competitor social media strategies. It processes data from Twitter/X, Reddit, LinkedIn, and other platforms to produce actionable intelligence about market sentiment, customer opinions, and industry trends. This skill is essential for market research, brand monitoring, and competitive analysis.

<https://github.com/anthropics/claude-skills>

Social Intelligence

Autoresearch Skill \$\$\$

An autonomous research agent that can be given a broad topic and will independently conduct a comprehensive research investigation over an extended period. Autoresearch Skill plans its own research strategy, identifies the most relevant sources, takes notes, connects findings across domains, and produces a final synthesis report. Unlike single-query research tools, Autoresearch works over minutes or hours, iteratively deepening its understanding. It's the closest thing to having a dedicated research analyst that works 24/7. This skill is ideal for comprehensive market research, technology landscape analysis, and due diligence investigations.

<https://github.com/anthropics/claude-skills>

Autonomous Research

DNA STRAND

Research DNA — the competitive advantage that compounds. Every research project makes you more knowledgeable, making the next project faster and more valuable. Over 12 months, a developer who consistently uses AI research tools builds a knowledge base that would be impossible to accumulate manually. This creates a "knowledge moat" — a depth of understanding in your domain that competitors simply cannot match regardless of their budget. The research DNA is what separates consultants who charge *100/hour* from those who charge *500/hour* for the same deliverable in less time.

EXERCISE

Run a Deep Research Sprint

1. **Choose a research topic:** Pick a topic relevant to your business, career, or a project you're considering. The topic should be specific enough to be researchable but broad enough to require multiple sources. Write your research question in one clear sentence.
2. **Plan your research strategy:** Before using any AI tools, write down: what you already know about this topic, what specific questions you need to answer, what types of sources would be most credible, and what the final output should look like (report, presentation, decision memo).
3. **Execute the research pipeline:** Use the Deep Research Engine or a combination of AI tools to conduct your research. Start with a broad search, identify 5-10 key sources, extract the most relevant findings from each, and note any contradictions or gaps in the available information.
4. **Synthesize your findings:** Compile your research into a structured report with the following sections: Executive Summary (200 words), Key Findings (3-5 bullet points with evidence), Detailed Analysis (500+ words), Sources Consulted (list with URLs), and Recommendations (3-5 actionable next steps).
5. **Apply the Evidence-Based Dialogue check:** Review your report and identify any claims that lack proper sourcing. Rate your confidence in each key finding on a scale of 1-5. Note what additional research would be needed to increase confidence in your lowest-rated findings.

Write your answer here:

QUIZ

1. What makes the Deep Research Engine different from a single AI search query?

- A) It's faster B) It decomposes questions, searches multiple sources, evaluates credibility, and iteratively refines its search C) It only uses academic sources D) It requires no internet connection

2. Which research skill acts as a quality control layer by verifying claims against reliable sources?

- A) Deep Research Engine B) Academic Research Skills C) Evidence-Based Dialogue D) Social Media Research

3. What is a "knowledge moat" in the context of AI research skills?

- A) A subscription to research databases B) A depth of domain understanding built over time that competitors cannot easily replicate C) A firewall around your data D) A collection of saved search queries

4. What time reduction can AI research tools achieve for a comprehensive market analysis?

- A) 10% faster B) 2 weeks reduced to 2 hours C) No real time savings, just better quality D) Only useful for academic research

5. The Autoresearch Skill differs from other research tools primarily because it:

- A) Uses only social media sources B) Works autonomously over extended periods, iteratively deepening its understanding C) Only searches one database D) Requires constant human guidance

MODULE 8

GIT MONEY — Product & Building

7 skills for turning ideas into revenue-generating products

Every skill you've learned in this workbook converges in this final module — product and building skills. This is where the DNA of technical capability meets the MONEY of business acumen. The skills in this module cover the complete product development lifecycle: from discovering what people actually need (Jobs To Be Done interviews, AI Transformation Discovery), to defining and planning the product (PM Skills Marketplace), to building the actual product (AI Video Toolkit, AI

Music Album), to understanding and distributing it (Repomix, Skill Seekers). The critical insight here is that the gap between "having AI skills" and "making money with AI skills" is the product layer. You can be the best Claude operator in the world, but if you can't identify market opportunities, define products, and ship them, your skills remain latent potential. Product skills are the activator that turns raw technical capability into revenue. The most successful AI-native entrepreneurs in 2025 share a common pattern: they identify a specific pain point, build a minimal product using AI tools in 1-4 weeks, launch it, gather feedback, and iterate rapidly. They don't spend months planning — they ship fast and learn from the market. The skills in this module give you the frameworks and tools to follow this pattern reliably.

PM Skills Marketplace \$\$\$

A curated marketplace and framework for product management skills specifically designed for AI-native products. PM Skills Marketplace provides templates, frameworks, and best practices for product discovery, requirements definition, roadmap planning, and feature prioritization in the context of AI-powered software. It includes AI-specific considerations like model selection, prompt engineering for product features, user feedback loops, and ethical AI deployment. This skill bridges the gap between technical AI capability and product thinking — essential for anyone building AI products.

<https://github.com/anthropics/claude-skills>

Product Management

JTBD Interview Tool \$\$\$

A Jobs To Be Done (JTBD) interview framework powered by AI that helps you discover what customers actually hire products to do. The JTBD Interview Tool guides you through conducting customer interviews that reveal the underlying jobs, pains, and gains that drive purchasing decisions. Instead of asking customers what features they want (which leads to building the wrong thing), JTBD interviews uncover the progress customers are trying to make in their lives. The AI helps design interview questions, analyze responses, identify patterns, and synthesize findings into actionable product insights. This is the single most valuable customer research skill for product builders.

<https://github.com/anthropics/claude-skills>

Customer Research

AI Transformation Discovery \$\$\$

A consulting framework for identifying and prioritizing AI transformation opportunities within organizations. AI Transformation Discovery helps you assess a company's current processes, identify where AI can deliver the most value, estimate implementation effort and ROI, and create a prioritized roadmap. The framework includes assessment templates, process mapping tools, and value estimation models. This skill is the foundation of an AI consulting practice — companies will pay \$5,000-50,000 for a thorough AI transformation assessment that identifies their highest-value automation opportunities.

<https://github.com/anthropics/claude-skills>

AI Consulting

AI Video Toolkit \$\$

A comprehensive toolkit for creating AI-generated and AI-assisted video content. The AI Video Toolkit covers the full video production pipeline: script writing, storyboarding, voice generation, visual asset creation, editing, and post-production. It integrates with tools for text-to-video generation, AI voice cloning, automated captioning, and visual effects. Video content is the fastest-growing content format, and AI makes professional-quality video production accessible to creators without production experience. This skill enables video production services, course creation, and content marketing at scale.

<https://github.com/anthropics/claude-skills>

Video Production

AI Music Album \$\$

A creative skill for producing complete AI-generated music albums — from concept to final tracks. The AI Music Album skill covers genre selection, lyric writing, melody composition, arrangement, and production. It leverages AI music generation tools to create original tracks that can be used for background music, podcasts, videos, games, and commercial projects. While AI-generated music won't replace human artistry, it fills a massive demand for royalty-free, custom music for content creators and businesses. This skill opens opportunities in music licensing, content creation, and audio branding.

<https://github.com/anthropics/claude-skills>

Music Production

Repomix

A tool that packs entire code repositories into a single, well-structured file optimized for LLM consumption. Repomix traverses a codebase, respects .gitignore rules, handles binary files appropriately, and produces a single text file that contains all the relevant code organized by directory structure. This is essential when you need to give an AI model complete context about a codebase — for code review, documentation generation, migration planning, or onboarding assistance. Repomix turns the complex task of codebase understanding into a simple file upload.

<https://github.com/yamadashy/repomix>

Codebase Tools

Skill Seekers

A discovery and evaluation framework for finding the most valuable AI skills and tools for specific use cases. Skill Seekers helps you navigate the overwhelming landscape of AI capabilities by matching your needs to the best available skills. It provides evaluation criteria, comparison matrices, and recommendation engines that consider factors like quality, reliability, cost, and community support. For builders and consultants, Skill Seekers ensures you're always using the best tool for each job, staying current with the rapidly evolving AI tool landscape.

<https://github.com/anthropics/claude-skills>

Tool Discovery

DNA STRAND

Product DNA — from skill to revenue. The gap between knowing AI tools and making money with them is the product layer. Every successful AI entrepreneur follows the same pattern: identify a pain point, build a minimal solution in 1-4 weeks using AI tools, launch it fast, learn from real users, and iterate. The skills in this module — from JTBD interviews that uncover real needs, to PM frameworks that define products, to AI Transformation Discovery that identifies consulting opportunities — are the bridge between technical ability and business impact. Your product DNA is what turns your GITBUILT skills into a sustainable income.

EXERCISE

Design Your AI Product

- 1. Pick a niche:** Choose a specific industry or profession where AI could solve a real problem. Write down the niche, why you chose it, and what initial research you've done (or plan to do) to validate the opportunity. Your niche should be specific enough to target but large enough to support a business.
- 2. Define the Jobs To Be Done:** Using the JTBD framework, identify 3 specific "jobs" that people in your niche are trying to accomplish. For each job, write: the job statement ("When [situation], I want to [motivation], so I can [outcome]"), the main pain points, and the current alternatives they're using (including doing nothing).
- 3. Sketch the MVP:** Based on your JTBD analysis, define the minimum viable product that addresses the most important job. Write a product brief that includes: product name, one-sentence description, 3 core features, the technology stack (which AI tools/models you'll use), and what success looks like after 30 days of launch.
- 4. Plan the launch:** Create a launch plan with the following elements: where you'll find your first 10 users, what you'll charge (or if free, how you'll monetize later), what metrics you'll track, and what your week-1, week-2, and week-4 milestones are. Be specific and realistic.
- 5. Identify your skill gaps:** Looking at the product you've designed, which skills from this workbook do you need to develop further? Which modules are most relevant? Create a personal learning plan for the next 30 days to close your biggest skill gaps.

Write your answer here:

QUIZ

1. What is the core principle behind the Jobs To Be Done (JTBD) framework?

- A) Ask customers what features they want B) Discover what progress customers are trying to make in their lives, not what features they request C) Build the cheapest possible product D) Copy what competitors are doing

2. What is the typical price range for an AI Transformation Discovery assessment?

- A) 100 – 500 B) 1,000-3,000 C) 5,000 – 50,000 D) 100,000+

3. What does Repomix do and why is it useful?

- A) It compiles code for deployment B) It packs entire codebases into a single file optimized for LLM consumption C) It generates unit tests D) It manages git repositories

4. What common pattern do the most successful AI-native entrepreneurs follow?

- A) Spend 6 months planning before building B) Raise venture capital first C) Identify a pain point, build an MVP in 1-4 weeks with AI tools, launch fast, and iterate based on feedback D) Only build for enterprise clients

5. Which product skill is described as "the single most valuable customer research skill for product builders"?

- A) PM Skills Marketplace B) AI Transformation Discovery C) JTBD Interview Tool D) Skill Seekers

Cheat Sheet

All 56 skills across 8 modules — quick reference table

#	Skill	Category	Tag	One-Liner
1	Claude Code	Coding Agent	DNA	Anthropic's terminal-native coding agent for your entire repo
2	Cursor	Agent IDE	DNA	AI-native IDE with multi-file Composer feature
3	Codex CLI	Coding Agent	DNA	OpenAI's open-source terminal coding agent
4	Windsurf	Agent IDE	DNA	Codeium's AI IDE with Cascade agentic workflows
5	Superpowers	Coding Skill	DNA	Enhanced Claude workflows for code review and refactoring

6	Spec Kit	Spec-Driven	DNA	Specification-first development with AI implementation
7	Aider	Pair Programming	DNA	Open-source AI pair programmer with deep Git integration
8	OpenClaw	Agent Framework	DNA	Open-source general-purpose agent framework
9	LangGraph	Graph Agent	DNA	Graph-based framework for stateful multi-step agents
10	CrewAI	Multi-Agent	DNA	Multi-agent orchestration with role-based collaboration
11	AutoGPT	Autonomous	DNA	Pioneering autonomous agent that sparked the agent revolution
12	Dify	Visual Platform	DNA	Visual workflow builder for AI apps and RAG pipelines
13	OWL	Lightweight	DNA	Minimal boilerplate framework for task-specific agents
14	CopilotKit	Copilot	DNA	React framework for adding AI copilots to web apps
15	pydantic-ai	Type-Safe	DNA	Type-safe agent framework using Pydantic models
16	Tavily MCP	MCP Search	DNA	Real-time web search optimized for AI agents
17	Context7	MCP Docs	DNA	Up-to-date library documentation on demand
18	Task Master AI	MCP Tasks	DNA	Persistent task management for AI coding sessions
19	MCP Playwright	MCP Browser	DNA	Browser automation through the Playwright library
20	fastmcp	MCP Builder	DNA	Python library for rapidly building custom MCP servers
21	markdownify-mcp	MCP Convert	DNA	Converts web pages to clean Markdown for LLMs
22	MCPHub	MCP Registry	DNA	Discovery platform and registry for MCP servers
23	Ollama	Local LLM	DNA	Run LLMs locally with a single command
24	Open WebUI	Local Interface	DNA	Self-hosted ChatGPT-like interface for local models
25	LlamaFile	Portable AI	DNA	Package LLMs into single executable files
26	Unsloth	Fine-Tuning	DNA	2-5x faster fine-tuning with 80% less memory
27	vLLM + FastAPI	Model Serving	DNA	Production-grade model serving with PagedAttention
28	Portkey Gateway	AI Gateway	DNA	Observability, caching, and load balancing for AI APIs
29	frontend-design	UI/UX	MONEY	Generate production-ready frontend UI from text descriptions
30	canvas-design	Visual Design	MONEY	Visual compositions for print and digital formats
31	algorithmic-art	Gen Art	MONEY	Unique visual artworks via mathematical algorithms
32	Color Expert	Branding	MONEY	Color theory, palette generation, and WCAG compliance
33	Hand-Drawn Diagrams	Diagrams	MONEY	Sketch-style diagrams for presentations and docs
34	claudedesignskills	3D/Motion	MONEY	3D rendering, motion graphics, and visual effects
35	Nothing Design	Minimalist	MONEY	Minimalist design through reduction and negative space
36	GPT Image 2	Image Gen	MONEY	Photorealistic image generation from text descriptions
37	Social Media OS	Content System	MONEY	Complete social media operating system for management

38	hook-generator	Content	MONEY	Attention-grabbing hooks for any social platform
39	post-scorer	Optimization	MONEY	Predict engagement performance before publishing
40	youtube-thumbnail	YouTube	MONEY	CTR-optimized thumbnails for YouTube videos
41	Email Marketing Bible	Email	MONEY	Full lifecycle email marketing strategies and automation
42	GEO/SEO Claude	SEO/GEO	MONEY	Dual optimization for search and AI answer engines
43	Twitter Algorithm Optimizer	Twitter/X	MONEY	Content optimization for Twitter/X algorithm
44	Competitive Ads Extractor	Intel	MONEY	Extract strategic insights from competitor ad campaigns
45	Deep Research Engine	Research	MONEY	Multi-step comprehensive research with citations
46	Academic Research Skills	Academic	MONEY	Academic literature search and synthesis
47	Evidence-Based Dialogue	Quality	MONEY	Verify AI claims against reliable, verifiable sources
48	Social Media Research	Social Intel	MONEY	Extract insights from social media conversations
49	Autoresearch Skill	Auto Research	MONEY	Autonomous research agent for extended investigations
50	PM Skills Marketplace	Product	MONEY	PM frameworks for AI-native product development
51	JTBD Interview Tool	Research	MONEY	Jobs To Be Done customer interview framework
52	AI Transformation Discovery	Consulting	MONEY	Identify and prioritize AI transformation opportunities
53	AI Video Toolkit	Video	MONEY	Full pipeline AI video production toolkit
54	AI Music Album	Music	MONEY	Complete AI-generated music album production
55	Repomix	Codebase	MONEY	Pack entire repos into LLM-optimized single files
56	Skill Seekers	Tool Discovery	MONEY	Discover and evaluate the best AI skills for any use case

Final Assessment

FINAL ASSESSMENT

1. What is Claude Code's primary interface and where does it operate?

- A) Web browser interface B) Mobile app C) Command-line terminal integrated with your repository D) Desktop GUI application
-

2. Which coding agent is built directly on VS Code and features a "Composer" for multi-file edits?

- A) Claude Code B) Windsurf C) Cursor D) Aider
-

3. In LangGraph, agent workflows are defined as what type of structure?

- A) Decision trees B) Neural networks C) Directed graphs with nodes and edges D) Linear pipelines
-

4. CrewAI's unique approach involves creating what type of system?

- A) Single powerful agent B) Teams of specialized agents with defined roles that collaborate C) Rule-based expert system D) Static workflow automation
-

5. What does MCP stand for and what is its primary purpose?

- A) Machine Control Protocol — for robotics B) Model Context Protocol — standard for connecting AI models to tools C) Memory Cache Protocol — for data storage D) Multi-Cloud Platform — for deployment
-

6. Which tool converts any LLM into a single portable executable file?

- A) Ollama B) vLLM C) Unsloth D) LlamaFile
-

7. What technique does vLLM use to achieve high-throughput model serving?

- A) Model distillation B) Quantization C) PagedAttention and continuous batching D) Knowledge distillation
-

8. Which design skill category commands the highest per-project rates (\$1,000-10,000)?

- A) Frontend design B) Color Expert C) 3D/Motion design D) Canvas design
-

FINAL ASSESSMENT

1. What is Claude Code's primary interface and where does it operate?

- A) Web browser interface B) Mobile app C) Command-line terminal integrated with your repository D) Desktop GUI application
-

2. Which coding agent is built directly on VS Code and features a "Composer" for multi-file edits?

- A) Claude Code B) Windsurf C) Cursor D) Aider
-

3. In LangGraph, agent workflows are defined as what type of structure?

- A) Decision trees B) Neural networks C) Directed graphs with nodes and edges D) Linear pipelines
-

4. CrewAI's unique approach involves creating what type of system?

- A) Single powerful agent B) Teams of specialized agents with defined roles that collaborate C) Rule-based expert system D) Static workflow automation
-

5. What does MCP stand for and what is its primary purpose?

- A) Machine Control Protocol — for robotics B) Model Context Protocol — standard for connecting AI models to tools C) Memory Cache Protocol — for data storage D) Multi-Cloud Platform — for deployment
-

6. Which tool converts any LLM into a single portable executable file?

- A) Ollama B) vLLM C) Unsloth D) LlamaFile
-

7. What technique does vLLM use to achieve high-throughput model serving?

- A) Model distillation B) Quantization C) PagedAttention and continuous batching D) Knowledge distillation
-

8. Which design skill category commands the highest per-project rates (\$1,000-10,000)?

- A) Frontend design B) Color Expert C) 3D/Motion design D) Canvas design
-

9. What is the estimated ROI for email marketing per \$1 spent?

- A) 5 B) 12 C) 36 D) 100
-

10. GEO (Generative Engine Optimization) optimizes content for what new type of search?

- A) Voice search B) Image search C) AI-powered answer engines like ChatGPT and Perplexity
D) Video search
-

11. What does the Evidence-Based Dialogue skill ensure about AI-generated research?

- A) It's written in formal language B) It's as long as possible C) Claims are verified against reliable sources with confidence levels D) It uses only academic sources
-

12. The JTBD Interview Tool helps you discover what fundamental insight about customers?

- A) Their budget B) What progress they're trying to make in their lives C) Their social media habits D) Their technical expertise
-

13. What is the typical price range for an AI Transformation Discovery consulting assessment?

- A) 100 – 500 B) 500-2,000 C) 5,000 – 50,000 D) 500,000+
-

14. Which framework provides a visual drag-and-drop builder for AI agent workflows?

- A) LangGraph B) pydantic-ai C) Dify D) OWL
-

15. What does Repomix do?

- A) Manages package dependencies B) Creates git repositories C) Packs entire codebases into a single LLM-optimized file D) Generates README files
-

16. A freelancer delivering a 20-hour project in 4 hours at \$100/hour effectively earns what rate?

- A) 200/hour B) 400/hour C) 500/hour D) 1,000/hour
-

17. What makes Unsloth significant for model fine-tuning?

- A) It fine-tunes without any data B) It runs 2-5x faster with 80% less memory, enabling fine-tuning on consumer GPUs C) It only works with small models D) It eliminates the need for training
-

